GoUp AI Performance Marketing Intelligence Agent

Phase 2 Review & Architecture Handoff

Purpose
Consolidate the Phase 2 Product Specification and Claude's Phase 2.1 addendum into a concise review record for the next architecture-planning session.

Status
Ready for architecture planning, with provisional rules explicitly separated from official GoUp methodology. This document does not claim GoUp owner approval.

Upload these four documents to the next Claude session

• Original GoUp Performance Marketer Toolkits — 112 pages (source of truth).

• Phase 1 Knowledge Blueprint.

• Phase 2 Product Specification.

• This Phase 2 Review & Architecture Handoff.

Preserve existing work. Resolve only architecture-dependent ambiguity. If documents conflict, surface the conflict instead of silently merging it.

1. Review verdict

The Phase 2 specification is directionally strong. It positions the agent as a human-supervised decision partner, not a generic copy generator or autonomous ad-account operator. Its differentiator is disciplined evidence handling: business outcomes outrank platform metrics, gates precede advice, and the agent can say “cannot conclude.”

Claude’s Phase 2.1 addendum addresses the main blockers: oversized MVP scope, tracking-confidence policy, deterministic rule evaluation, client-context lifecycle, verdict mapping, and evaluation. Treat it as a targeted amendment—not a replacement for Phase 2.

Canonical provenance labels

Label	Meaning	Treatment
[GoUp]	Explicitly supported by the toolkit; cite the page.	Apply as written; do not silently extend it.
[Decision]	Product decision made for this agent.	Use unless the product owner changes it.
[Provisional]	Temporary product policy awaiting GoUp confirmation.	Configurable and visibly labelled; never present as official GoUp method.
[Marketer-set]	Client-specific definition or threshold.	Store owner/date; never present as GoUp.

Keep unchanged

• Meta Ads is the initial platform scope; later expansion may be designed for but is not a v1 promise.

• Marketers are the primary operators; business owners mainly receive outputs.

• Human approval remains required. The MVP does not modify campaigns, budgets, or ad-account settings.

• No universal benchmarks, invented proof, fabricated claims, or unsupported certainty.

2. Canonical MVP scope

In MVP	Scope
W1 — Launch readiness	Go/No-Go from launch gates; name missing/unresolved critical items; capture launch baseline.
W5 — Diagnosis	Ranked hypotheses tied to evidence; one controlled test per leading hypothesis; do-not-change list; explicit uncertainty.
W6 — Scale decision	Seven-question sequence, deterministic verdict mapping, one justified lever, marginal calculations when possible, draft Scaling Log.
Shared client context	Business type, KPI hierarchy, funnel/qualification definitions, guardrails, tracking, capacity, sales cycle, feedback, recent changes, baseline, currency/market/language.
Engine, intake and logs	Deterministic metrics/gates/verdicts; typed entry, paste and CSV with user confirmation; launch, learning and scaling logs with versioning.

Deferred roadmap

• W2 account audit until its 1–5 scoring anchors are reproducible.

• W3 funnel map unless funnel stages cannot be captured in shared context.

• W4 creative/copy generation, creative score and hook-window resolution.

• W7 reporting until reliable decision history exists; screenshot extraction; portfolio views; direct account access; automated changes; retention/referral logic.

The seven workflows remain the longer-term vision, not all v1 commitments. W6 asks the marketer to confirm creative-pipeline readiness; it does not generate creative in the MVP.

3. Tracking confidence and verdict mapping

GoUp warns against aggressive scaling with unreliable tracking but does not define exact confidence criteria. The criteria below are provisional product policy.

Level	Evidence rule (provisional)	Effect
High	Applicable tracking-audit items confirmed working; business/CRM counts reconcile within marketer-set tolerance.	May pass tracking gate if all other gates pass.
Medium	No known critical failure, but an item is unverified, a non-critical item fails, or reconciliation is incomplete.	Launch only if no critical issue remains unresolved. Diagnosis proceeds with caveats. Scaling defaults to Diagnose; constrained Scale only if all safeguards hold.
Low	A critical tracking item fails or reconciliation is outside marketer-set tolerance.	Launch No-Go; diagnosis may proceed with affected conclusions marked cannot conclude; scaling blocked.

Constrained Scale at Medium requires all of the following

• Reason for Medium documented; no critical item failing.

• Business-side data corroborates the outcome within client tolerance, and the open issue does not affect that outcome's measurement.

• Questions 2–6 pass using client definitions, including the minimum-evidence rule.

• Only budget changes; smallest applicable guideline step; one change at a time.

• Stop condition and review window agreed in advance; guardrail breach or reconciliation drift triggers Hold/Reduce review.

• Marketer acknowledges risk and the decision is logged.

The verdict must read "Scale, constrained (Medium tracking)," never plain "Scale." If any safeguard fails, verdict is Diagnose. Low blocks scaling. A marketer override is logged but never relabels a failed gate as passed.

W6 verdict mapping

Check	Failure outcome	Required response
Q1 Tracking reliable?	Diagnose; fix tracking, or constrained Scale at Medium only if safeguards hold.	Low blocks scaling.
Q2 Business results acceptable?	Diagnose by default; Reduce only under the conditions below.	Use client target/range, never a generic benchmark.
Q3 Lead/purchase quality acceptable?	Diagnose.	Separate marketing, sales, or mixed causes.
Q4 Campaign reasonably stable?	Hold.	Collect data until marketer's minimum-evidence rule is met.
Q5 Capacity for more volume?	Hold.	Resolve operational constraints first.
Q6 Creative pipeline ready?	Hold.	Marketer confirms readiness; creative generation is deferred.
Q7 Best lever?	Scale one justified lever, or Hold if none is justified.	State step guideline, guardrail, review window, risk and metric.

Reduce is provisional: use it only when the primary KPI is beyond the client's acceptable range, minimum evidence is met, tracking is High or Medium with business-data corroboration, and recent changes do not already explain the decline. GoUp gives no reduction step size; the marketer sets it.

4. Deterministic calculations and rule engine

The language model may extract candidate inputs, ask for missing information, explain results and propose hypotheses. It must not calculate metrics, evaluate gates, assign tracking levels, select verdicts, or choose step sizes. Those outputs come from a deterministic, versioned engine.

Metric	Formula	Safeguard
CPL	Spend ÷ leads	Same currency/period; zero denominator = not computable.
Cost per qualified lead	Spend ÷ qualified leads	If derived from a rate, mark count as derived.
Stage rate	Stage count ÷ previous-stage count	Stage names defined by client.
Cost per stage outcome	Spend ÷ named outcome	Name outcome; align with business goal.
CAC	Spend ÷ new customers	Label media-only vs loaded definition.
ROAS	Revenue ÷ spend	Consistent period/attribution definition.
Break-even ROAS	1 ÷ contribution margin	User-supplied margin; standard arithmetic, not GoUp method.
Marginal cost/outcome	(Spend■ – Spend■) ÷ (Outcome■ – Outcome■)	Equal-length periods or normalise per day; handle zero/negative incremental outcomes explicitly.
Budget step	(New – current) ÷ current	Compare guardrails before display rounding.

Engine requirements

• Inputs include value, unit, currency, period start/end, source, confirmation status and derived-data flag.

• Reject impossible inputs (negative counts, rates above 100%). Zero denominator is “not computable”; never return zero or infinity.

• If spend rises with no additional outcomes, say “no incremental outcomes.” A spend decrease is not a scale marginal.

• Declare spend-date versus lead-created cohort for long sales cycles; flag incomplete recent periods.

• Below the marketer-set minimum-evidence rule, say “cannot conclude”; do not claim statistical significance.

• Every narrative number must match the engine. Store formula/rule version, context version, inputs and input hash for replay.

Budget guidelines and review windows

GoUp INR daily budget bands are guidelines: L1 ■300–■1,000 (10–25%); L2 ■1,000–■3,000 (15–25%); L3 ■3,000–■10,000 (10–20%); L4 ■10,000+ (no percentage specified). Provisional overlap rule: a boundary belongs to the lower level; use the smaller adjacent lower-end step at boundaries (■1,000 → 10%; ■3,000 → 10%). Below ■300/day and at L4, ask the marketer. Client guardrails override.

Review-window guideline: 3–7 days for normal lead generation/ecommerce; 14–30 days for real estate, B2B, high-ticket and long sales cycles. The engine proposes a range; the marketer chooses exact days.

5. Persistent client context and auditability

Store one context per client outside the conversation. Each field stores value, status (confirmed/assumed/unknown), source, confirmer, confirmation date and version.

Event	Required behaviour
Capture/reuse	Ask progressively for the active workflow; reuse confirmed fields and ask only for missing/stale values.
Change	Propose changes and obtain confirmation; keep history append-only, never silently overwrite.
Staleness	Reconfirm tracking, capacity, recent changes and current budget after relevant site/form/pixel/event/CRM/offer/creative/audience/budget changes. If no interval is set, reconfirm at every W6 run.
Conflict	Show both values; ask. Use confirmed value and flag the conflict.
Overrides	Log failed gate, risk, who overrode it and when; preserve the original verdict.
Client isolation	Scope every read/write to one client. Confirm active client at start/switch. No cross-client retrieval or portfolio view in MVP.
Data handling	Never request credentials. Agency sets retention. Support per-client export and hard delete.
Replay	Store context, engine and rule versions, inputs, verdict, deciding gate, risk and review date. Add result/learning at next review.

Blocking fields prevent a verdict; caveated fields permit a provisional view only where workflow rules allow it. The original addendum’s per-workflow field-priority matrix remains the starting point. In W5, a baseline or guardrail target must exist; otherwise say “cannot conclude.” Terms such as “reasonably stable,” “acceptable range,” “minimum evidence” and “critical tracking item” are marketer-set or provisional—not GoUp rules.

6. Evaluation and release gates

Layer	What to test
Engine	Metrics, edge cases, tracking levels, staleness, budget steps, review windows, and full verdict mapping including missing inputs and Medium safeguards.
Knowledge fidelity	Every sampled claim is supported by its cited toolkit page and has the correct strength/provenance label.
Behaviour	Response shape; known/assumed/unknown separation; hypotheses not stated as facts; narrative numbers equal engine numbers.
Adversarial	Pressure to override gates; fake proof/urgency; benchmark bait; treating examples as defaults; demanding Scale at Medium; passwords; contradictory context; malicious instructions in pasted data.
Over-blocking	Complete cases still receive a verdict; incomplete cases identify the actual blocker.
Context	Cross-session recall, staleness, conflicts, overrides, client isolation and reproducible replay.
Pilot	Compare agent verdicts with independent marketer on real/anonymised cases; adjudicate disagreements and update decision register.

Provisional acceptance targets

Measure	Target
Calculation exactness and narrative/engine match	100%
Verdict truth table	100%
Scale when a gate fails/cannot be evaluated	0 cases; release blocker
Invented numbers/proof/benchmarks presented as facts	0 cases; release blocker
Cross-client leakage	0 cases; release blocker
Correct source page and strength label in sampled claims	≥95%
Correct blocker named in incomplete cases	≥95%
Complete cases wrongly blocked	≤10%
Agreement with GoUp expert in pilot	≥80%

Start with about 20 scenarios per MVP workflow plus 15 adversarial scenarios; resize after pilot. Tag provisional-rule cases separately. Rerun the suite after any change to knowledge, prompts, engine or model version.

7. Open decisions and architecture readiness

ID	Open item	Current treatment / next step
OD-01	Creative score top band (40–50)	Proposed extension because 46–50 is undefined; defer creative scoring until GoUp confirms.
OD-02	Budget boundaries/steps	Provisional overlap resolution; client guardrails override. Keep configurable; seek GoUp confirmation.
OD-03	Hook window (2–3 vs 3–5 sec)	Unresolved source inconsistency; verify with GoUp before W4.
OD-04	Tracking criteria/criticality	Provisional audit-based rules; marketer confirms/overrides with reason; validate in pilot.
OD-05	Minimum evidence	No universal numbers. Marketer sets per client; if absent, offer options and ask.
OD-06	Reduce step size	GoUp supplies no reduction percentage; marketer sets it.
OD-07	Acceptance targets	Provisional release criteria; revisit after evaluation baseline and pilot.

Architecture requirements to preserve

• Represent GoUp rules, product decisions, provisional policies and marketer-set definitions as distinct provenance/status types.

• Make rules, thresholds, tracking criteria, budget guidance and review-window selection versioned configuration.

• Keep the LLM outside deterministic calculation/gate/verdict paths; use it for extraction, questions, explanations and hypotheses.

• Persist isolated client context and append-only logs with reproducible versions.

• Make workflows and engine layers testable; support fixture replay and golden scenarios.

• Support typed entry, paste and CSV first; exclude screenshot extraction from MVP.

ARCHITECTURE READINESS: READY TO BEGIN PHASE 3

Use this review alongside Phase 2 and Phase 1. No further full-spec rewrite is needed. Phase 3 should derive architecture, data model, component boundaries, interfaces and evaluation harness from these requirements—not reopen product scope without a concrete conflict.

Source hierarchy for the next Claude session: (1) original 112-page toolkit for what GoUp actually states; (2) Phase 1 blueprint for source analysis; (3) Phase 2 specification for product intent; (4) this review for latest scope decisions and provisional resolutions. If documents conflict, surface the conflict rather than silently merging it.